

Parallel Programming

Brief overview

MPI(Message Passing Interface)

- MPI is a specification for the developers and users of message passing libraries.
- Portability - There is no need to modify your source code when you port your application to a different platform that supports (and is compliant with) the MPI standard
- In practice, MPI is a set of functions or subroutines used for exchanging data between processes.
- MPI has around 6 basic functions, and around 440+ [in MPI-3] functions/routines in total.
- Interface specifications have been defined for C/C++ and Fortran programs.

MPI Function/Subroutine

- **MPI_Init ()**: is the 1st function to call, initialize MPI variables. Creates MPI_COMM_WORLD communicator, which is list of all the connections among nodes.
- **MPI_Finalize()**: This is the last function to call which cleans up any job left by MPI.


Example:

```
#include "mpi.h"                // INCLUDE "mpif.h" for fortran
#include <stdio.h>                // part of regular c program
void main(argc, argv)
{
    MPI_Init( &argc, &argv );    // Starts MPI
    .....
    .....
    MPI_Finalize();              // Ends MPI
}
```


MPI Function/Subroutine (Contd.)

- `MPI_Comm_size(MPI_COMM_WORLD, nprocs, ierror)` //ierror is not in C, but in fortran
 - returns the number of nodes that the code is running on, as specified in the mpirun command (eg **mpirun** -np 8 myprog | tee output.txt)
- `MPI_Comm_rank (MPI_COMM_WORLD, myproc, ierror)`: returns node number, range 0 to (np-1).
- `MPI_Send (...)` and `MPI_Recv(...)`: For sending and receiving messages.
- **MPI_Bcast** (...) : broadcasts data from one node to all others. (Similar, **MPI_Scatter**)
- **MPI_Reduce** (...): Reverse of broadcast, all processors send to a single processor (similar, **MPI_Gather**).
- `MPI_Barrier (MPI_COMM_WORLD, ierror)`: force synchronization.


Bcast, Scatter, Gather, Reduction




broadcast



scatter



gather



reduction

MPI_Bcast

Broadcasts a message to all other processes of that group

count = 1;

source = 1;

broadcast originates in task 1





MPI_Bcast(&msg, count, MPI_INT, source, MPI_COMM_WORLD);

task 0





task 1

task 2

task 3



← msg (before)



← msg (after)

MPI_Scatter

Sends data from one task to all other tasks in a group


```
sendcnt = 1;
```

```
recvcnt = 1;
```

```
src = 1;
```

task 1 contains the message to be scattered

```
MPI_Scatter(sendbuf, sendcnt, MPI_INT,  
            recvbuf, recvcnt, MPI_INT,  
            src, MPI_COMM_WORLD);
```



MPI_Reduce

Perform and associate reduction operation across all tasks in the group and place the result in one task

count = 1;

dest = 1;

result will be placed in task 1

MPI_Reduce(sendbuf, recvbuf, count, MPI_INT, MPI_SUM,
dest, MPI_COMM_WORLD);

task 0

task 1

task 2

task 3

1

2

3

4

← sendbuf (before)

10

← recvbuf (after)

MPI_Gather

Gathers together values from a group of processes

```
sendcnt = 1;
```


```
recvcnt = 1;
```

```
src = 1;
```


messages will be gathered in task 1

```
MPI_Gather(sendbuf, sendcnt, MPI_INT,  
          recvbuf, recvcnt, MPI_INT,  
          src, MPI_COMM_WORLD);
```


task 0




task 1





task 2



task 3



← sendbuf (before)



← recvbuf (after)

MPI Data Type

Be familiar with MPI data type. Example:

MPI_CHAR

MPI_BYTE

MPI_SHORT

MPI_INT

..... <more>

Compiling ...

`mpicc -o hello_world hello-world.c`

Or, `mpicc -o hello_world hello-world.c -lm` // to link math library

`-mpicc //c`

`-mpif90 //fortran 90`

`-mpiCC // C++`

`-mpif77 //fortran 77`

Running...

`mpirun` -np 8 hello_world

Internet Resources

There are many sites, few are mentioned here:

For MPI:

<https://computing.llnl.gov/tutorials/mpi/>

<http://www.mcs.anl.gov/research/projects/mpi/tutorial/gropp/talk.html>

And <http://www.mcs.anl.gov/research/projects/mpi/tutorial/>

For Pthreads or, POSIX Thread:

<https://computing.llnl.gov/tutorials/pthreads/>

<http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>

+ example(s) in the text book.

Example: mpi_hello.c

```
#include <mpi.h>
#include <stdio.h>
#include <math.h>

int main (int argc, char* argv[])
{
    int rank, size, i, j;
    MPI_Init (&argc, &argv);          /* Initialization step to start MPI */
    MPI_Comm_rank (MPI_COMM_WORLD, &rank); /* get the id of the current process */
    MPI_Comm_size (MPI_COMM_WORLD, &size); /* get the number of processes involved */

    srand((int)time(NULL));             /* calling random number(RN) */
    int r = rand() % 20000;

    double y = r* sin((size-rank+1)*r); /* using sin to vary RN arbitrarily */
    r = abs(y);

    for (i=0; i<=r; i++)                /* To make random delay */
    {
        j=i*i*i;
    }

    printf( "Hello world from process %d of %d\n", rank, size);

    MPI_Finalize();                     /* Terminates MPI execution environment */

    return 0;
}
```