

# Test 1

Mengxiang Jiang  
CSEN 5322 Operating Systems

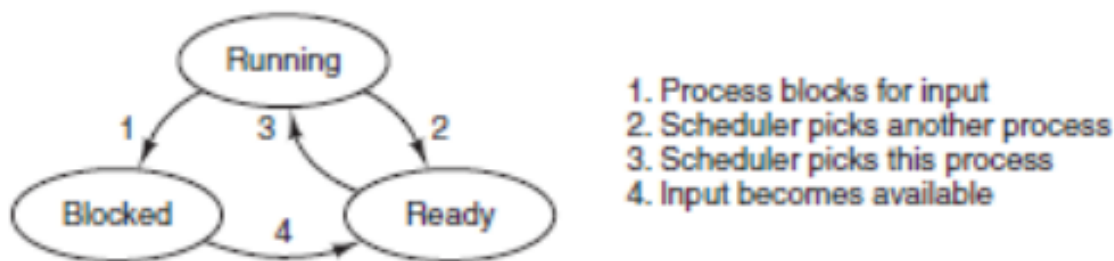
September 30, 2022

**Problem 1.** (a) What is multiprogramming? (b) What is the key difference between a trap and an interrupt?

(a) Multiprogramming is the ability for a user to run multiple programs at once.

(b) A trap is an instruction to switch the execution mode from user to kernel mode. An interrupt is a signal sent to the CPU usually when an I/O process has finished.

**Problem 2.** In Figure below, three process states and four transitions are shown. What are the circumstances in which either or both of the missing transitions might occur?

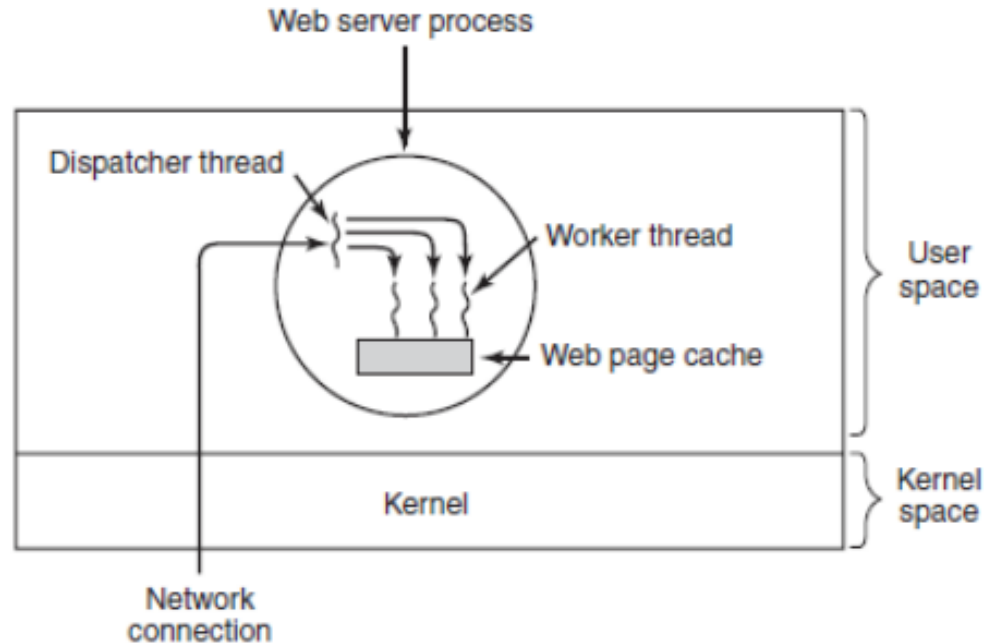


**Figure 1:** States of process

One missing transition is from Blocked to Running. This cannot happen since Blocked means it is currently unable to run until some external event happens, and when that something happens, it will go to Ready rather than Running. Of course, if there are no processes/threads in the Ready state, it will get immediately picked up and go to Running state, as if it transitioned to Running directly, but it still must go to Ready state first.

The second missing transition is from Ready to Blocked. This cannot happen since Ready means it is runnable but not currently running, so it cannot possibly be determined to be blocked on some external event. Of course, if it is picked up by Running and immediately requires I/O or some process that will block it, it will seem like it transitioned from Ready to Blocked, even though in actuality it first transitioned to Running first.

**Problem 3.** In Figure below, a multithreaded Web server is shown. If the only way to read from a file is the normal blocking read system call, do you think user-level threads or kernel-level threads are being used for the Web server? Why?



**Figure:** A multithreaded Web server.

Both user-level and kernel-level threads are used, since the user-level threads can read the local cache of a file if it exists, and if its not in the cache, then a kernel level thread will be used to fetch the file from disk.

**Problem 4.** (a) The aging algorithm with  $a = 1/2$  is being used to predict run times. The previous four runs, from oldest to most recent, are 80, 40, 80, and 30 msec. What are the predictions of the next time? (you must show your calculations)

Prediction 1 is just 80.

Prediction 2 is  $\frac{1}{2} * (80) + \frac{1}{2} * 40 = 60$

Prediction 3 is  $\frac{1}{2} * (60) + \frac{1}{2} * 80 = 70$

Prediction 4 is  $\frac{1}{2} * (70) + \frac{1}{2} * 30 = 50$

Therefore the sequence of predictions is 80, 60, 70, 50, so 50 is the next prediction.

(b) If a system has only two processes, does it make sense to use a barrier to synchronize them? Why or why not?

Yes, if you need to synchronize the two processes when one is faster than the other, a barrier is the ideal solution since it allows blocking without busy waiting (wasting cpu).

Although strict alternation can also work and is simpler to implement, it will waste cpu by busy waiting.

**Problem 5.** Suppose that we have a message-passing system using mailboxes. When sending to a full mailbox or trying to receive from an empty one, a process does not block. Instead, it gets an error code back. The process responds to the error code by just trying again, over and over, until it succeeds. Does this scheme lead to race conditions?

No race condition, since no shared variable is incorrectly read as a result.

**Problem 6.** Can two threads in the same process synchronize using a kernel semaphore if the threads are implemented by the kernel? What if they are implemented in user space? Assume that no threads in any other processes have access to the semaphore. Discuss your answers.

Two threads cannot synchronize using a kernel semaphore since the kernel does not know where the two threads come from and treats threads from different processes the same way it treats threads from the same process. They can synchronize if they are implemented in user space.

**Problem 7.** Multiple jobs can run in parallel and finish faster than if they had run sequentially. Suppose that two jobs, each of which needs 12 minutes of CPU time, start simultaneously. Assume 27% I/O wait is involved for both the jobs. (a) How long will the last one take to complete if they run sequentially? (b) How long if they run in parallel? You may leave the calculations without obtaining in the final single number format.

(a)  $\frac{12}{1-0.27} + \frac{12}{1-0.27} = 32.87$  minutes

(b)  $\frac{24}{(1-0.27^2)} = 25.89$  minutes